

BlueMountain

Enabling Automated, Rich, and Versatile Data
Management for Android Apps

*Sharath Chandrashekhara, Kyle Marcus, Rakesh G. M. Subramanya,
Hrishikesh S. Karve, Karthik Dantu and Steven Y. Ko*

Reliable Mobile Systems Lab
<http://www.nsr.cse.buffalo.edu>

Mobile Apps - State of Art

- Use local and cloud storage; rich forms of interaction, backup, sharing etc.
- Large companies use their own cloud; smaller developers use public cloud
- Too many cloud providers, no standard interface



Life as a Developer

- Several design choices
- Consistency models, interface and semantics
- Tangential and repetitive
- Binds an app to a particular cloud

⇒ *Developers want to reduce development time and provide more flexibility to users*




Life as a User

- Constrained by the developer's decisions
- Privacy concerns when data is moved to cloud
- Has to contact the developers for any customization


⇒ Users want flexibility and control



BlueMountain: Backup Scenario



BlueMountain: P2P Sharing Scenario



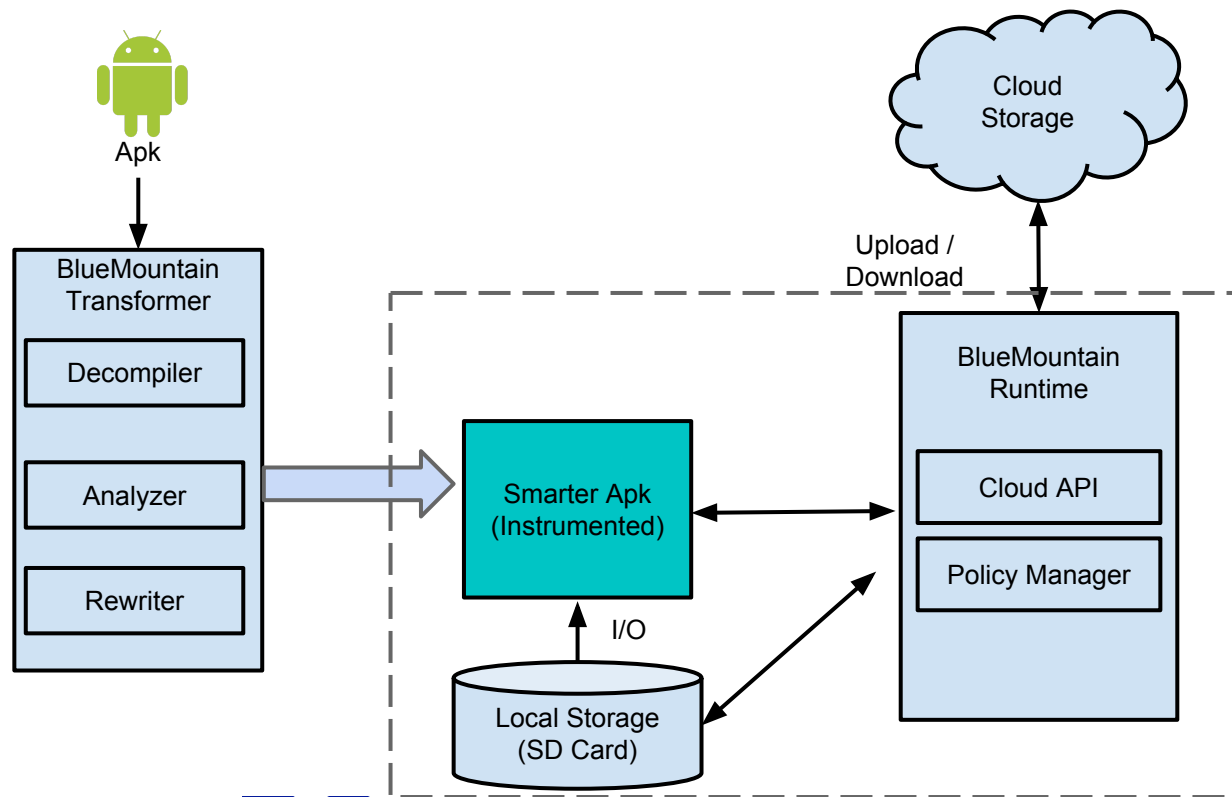
BlueMountain Goals

- **Reduce development effort:** Focus on app logic; treat all storage operations as local
- **Automatically transform apps:** Enable richer forms of data interaction
- **Flexibility:** Customize without access to source code
- **Post-development cycle:** Works with existing apps; no modifications to the Android platform for ease of deployment.


BlueMountain

- A system that automatically integrates cloud storage services with Android apps
- Main components
 - **App Transformer:** Analyses and rewrites app binaries by virtualizing the storage calls and enables richer data interactions
 - **Runtime:** Runs as a regular app; and interacts with the cloud and manages policies

BlueMountain Architecture



Challenges: Storage Virtualization



- Can we virtualize storage calls?
- Android options:
 - Files
 - Database
 - Key/Value

Storage Call Virtualization


```
public class MyFileOutputStream extends
FileOutputStream {
    public int write (Bytes b) {
        //Overriding
    }
}

public class main {
    public static void main (String args[]) {
        Bytes b = 10;
        MyFileOutputStream obj = new
        MyFileOutputStream ();
        myWrite (obj, b);
    }

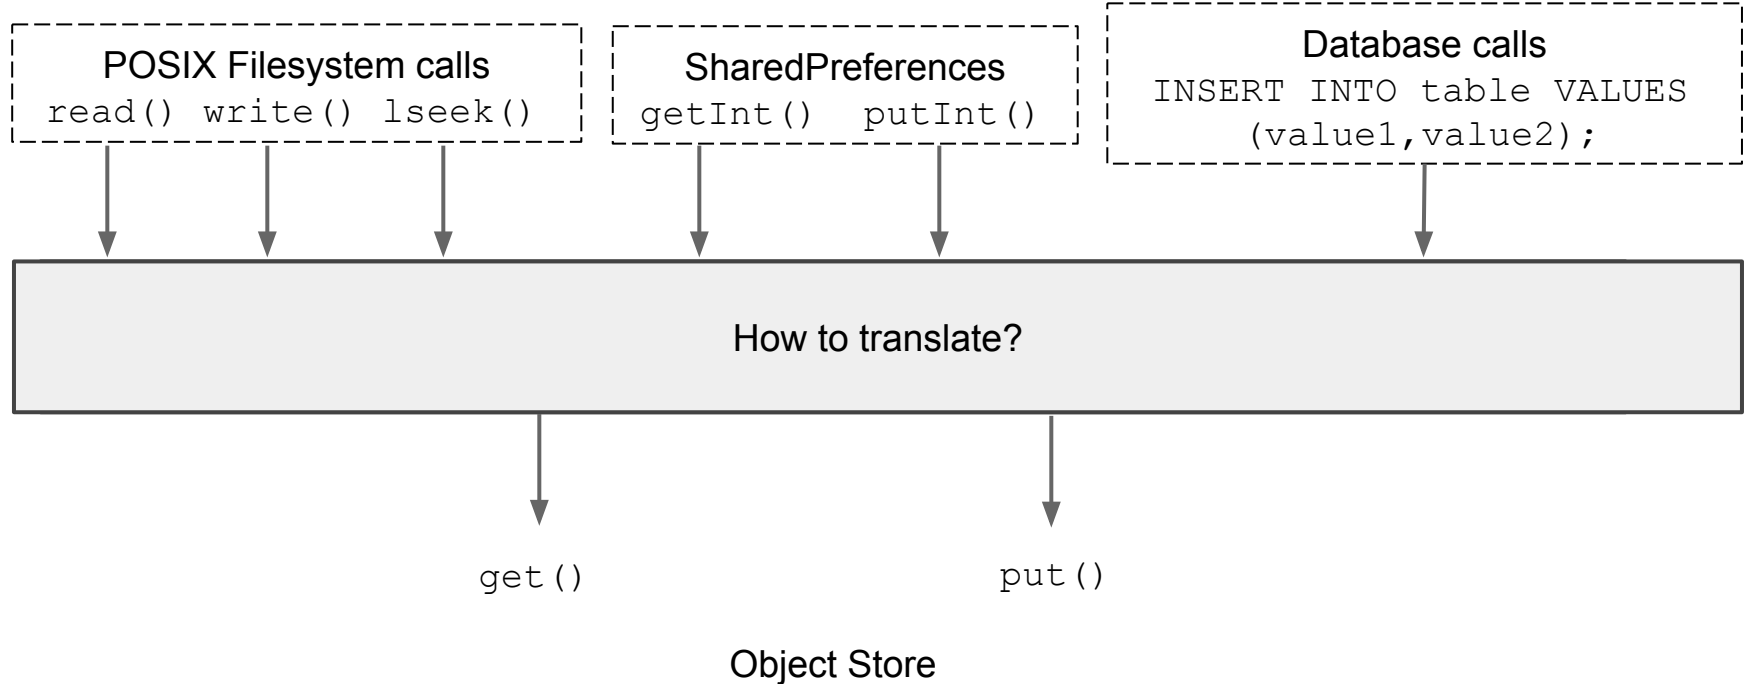
    public static void myWrite
    (FileOutputStream obj, Bytes b) {
        obj.write (b);
    }
}
```

- Need to **statically identify** all possible storage options and their APIs
- **More challenging than search and replace** because of polymorphism

Challenges: Cloud APIs



Challenges: Interface



Challenges: Semantics

- Handling concurrent updates
- Most clouds provide only eventual consistency
- Timing differences between local and cloud
- Time-bound eventually-consistent model?
- Getting additional inputs from the developers?

Related Work


- **Viewbox, Simba** - Fault tolerance and consistency guarantees
- **Cimbiosys** - Selective sharing of files
- **Procrastinator** - Rewriting the binary
- **MetaSync** - Enhances cloud services
- **CloudRail** - Unified API

Conclusion & Future Work

- Initial vision for BlueMountain - *storage virtualization* and *cloud integration*
- Developed a prototype to demonstrate these features
- Future work
 - Full implementation
 - Resolving interface and semantic mismatches
 - Analysing and categorizing the apps of Play Store and corporate apps which store data on their private cloud

BlueMountain

Enabling Automated, Rich, and Versatile Data Management for Android Apps



*Sharath Chandrashekhara, Kyle Marcus, Rakesh G. M. Subramanya,
Hrishikesh S. Karve, Karthik Dantu and Steven Y. Ko*
{sc296, kmarcus2, rakeshgu, hkarve, kdantu, stevko}@buffalo.edu
Reliable Mobile Systems Lab
<http://www.nsr.cse.buffalo.edu>